

THE DESK THAT GRADES ITSELF.

An options desk where **deterministic code builds every trade** and each analyst carries a track record. Analysts state probabilities, not verdicts, are Brier-scored against what actually happened, and one that is confidently wrong **loses its vote**. The models choose between pre-built candidates or refuse; they cannot invent a strike, because no code path allows it.

900

TESTS, NO NETWORK

3,686

CONTRACTS READ

632

STRUCTURES BUILT

1

CHOSEN, OR NONE

WHAT ONLY THIS DESK DOES

- › **Grades its own analysts.** Brier scores from resolved outcomes set voting weights each cycle. Unproven analysts stay at 1.0; a demoted one floors at 0.2 so it always has a path back.
- › **Runs an adversary** whose objections are journalled whether or not they prevail. In a recorded cycle its objection to a thin hedge moved the trader off the highest-credit candidate.
- › **Two vetoes that fail differently.** One is pure code checking position delta against the structure's own thesis; the other is a model shown the trade but never the committee's reasoning.
- › **Refusal is the product.** Two of four replayable scenarios are refusals.

ONE REAL DECISION, END TO END

snapshot	SPY 769.35, IV 2.27pp over realized
vol_analyst	p=0.63 • premium is rich
bear_adversary	p=0.40 • breakevens only 0.4% away
trader	picks c3, the wider cushion, answering the adversary
thesis veto	PASS • net delta -8.46 matches the thesis
blind veto	PASS • independent, never saw the debate
guard	ALLOW • risk \$322, 1 of 3 positions

The adversary changed the trade. That is the difference between a committee and a rubber stamp.

IS IT REASONING, OR REMEMBERING?

The standing criticism of LLM trading agents is **knowledge contamination**: papers such as TradingAgents (arXiv 2412.20138) backtest over dates inside the model's own training data, so a good result may be recall.

Our answer is checkable. The model's cutoff is **May 2026**; live decisions run on **August 2026** data. There was nothing to memorise. The walk-forward harness does span earlier dates, and that is acceptable only because it runs no model at all.

EVIDENCE, AND ITS LIMITS

SYMBOL	TRADES	WIN	EXPECTANCY	MAX DD
SPY	8	87.5%	+0.62R	2.0R
QQQ	7	85.7%	+0.57R	2.0R
AAPL	8	62.5%	-0.12R	6.0R
MSFT	7	85.7%	+0.57R	2.0R

Read this before the numbers. Thirty out-of-sample trades proves nothing statistically, and one symbol loses money, which is the point: an earlier version of this harness scaled its risk threshold to the wrong horizon and produced a 97% win rate by construction. It was caught and corrected before publication.

WHAT WE CANNOT PROVE

- › Paper trading only. No slippage or fill-quality evidence.
- › Few closed live trades, so most analyst weights are still 1.0. Demoting on a handful of outcomes is the error we refuse to make.
- › The fail-closed scenario is constructed, and its fixture says so.

VERIFY IT YOURSELF

- › `scripts/replay.py --all --verify` reproduces four recorded verdicts offline, with no credentials.
- › `make verify-journal` checks the hash chain.
- › Orders route through the official Alpaca CLI; market data through the official MCP server with **no trading toolset exposed**.